

System Design Interview Mastery Guide 2026

Complete preparation guide for Big Tech interviews

1. Complete Syllabus & Core Topics

Fundamentals

Scalability

- Horizontal vs Vertical scaling
- Auto-scaling strategies
- Stateless vs Stateful services
- Performance metrics (RPS, latency, throughput)

Reliability & Availability

- 99.9% vs 99.99% availability (downtime implications)
- Fault tolerance patterns: Circuit breakers, bulkheads, timeouts
- Disaster recovery and backup strategies
- Health checks and failover mechanisms

Consistency & CAP Theorem

- Strong vs Eventual consistency
- ACID properties in distributed systems
- BASE (Basically Available, Soft state, Eventual consistency)
- Partition tolerance trade-offs

System Architecture Patterns

Microservices Architecture

- Service decomposition strategies
- Inter-service communication (sync vs async)
- Service mesh (Istio, Linkerd)
- Domain-driven design principles

Serverless Computing

- Function-as-a-Service (FaaS) patterns
- Event-driven triggers

- Cold start optimization
- Cost implications and use cases

Event-Driven Architecture

- Event sourcing patterns
- Command Query Responsibility Segregation (CQRS)
- Saga pattern for distributed transactions
- Event streaming vs message queues

Database Systems

SQL Databases

- ACID compliance and isolation levels
- Indexing strategies (B-tree, Hash, Bitmap)
- Query optimization techniques
- Master-slave vs master-master replication

NoSQL Databases

- Document stores (MongoDB): Schema flexibility, horizontal scaling
- Key-value stores (Redis, DynamoDB): High performance, simple queries
- Column-family (Cassandra): Time-series data, write-heavy workloads
- Graph databases (Neo4j): Relationship-heavy data

Database Sharding & Partitioning

- Horizontal partitioning strategies
- Consistent hashing algorithms
- Cross-shard queries and joins
- Rebalancing and hotspot mitigation

Caching Strategies

Multi-Level Caching

- Browser cache (HTTP headers, service workers)
- CDN caching (geographic distribution)
- Application-level caching (in-memory, Redis)
- Database query result caching

Cache Patterns

- Cache-aside (lazy loading)
- Write-through and write-behind
- Cache invalidation strategies
- Cache warming and preloading

Load Balancing

Layer 4 vs Layer 7

- Network layer routing (IP, port)
- Application layer routing (HTTP headers, content)
- SSL termination considerations
- Performance implications

Load Balancing Algorithms

- Round robin and weighted round robin
- Least connections and least response time
- Consistent hashing for sticky sessions
- Health check strategies

Message Queues & Streaming

Apache Kafka

- Topic partitioning and consumer groups
- Exactly-once delivery semantics
- Kafka Connect for data integration
- Stream processing with Kafka Streams

Message Queue Patterns

- Point-to-point vs publish-subscribe
- Dead letter queues and retry mechanisms
- Message ordering guarantees
- Backpressure handling

API Design

REST API Best Practices

- Resource-based URL design

- HTTP status codes and error handling
- Pagination strategies (offset, cursor-based)
- API versioning (URL, header, content negotiation)

GraphQL Benefits

- Single endpoint, flexible queries
- Type system and schema validation
- Real-time subscriptions
- Caching challenges and solutions

gRPC for Internal Services

- Protocol Buffers serialization
- Bidirectional streaming
- Service discovery integration
- Performance benefits over REST

Security

Authentication & Authorization

- OAuth 2.0 flows and JWT tokens
- Role-based access control (RBAC)
- API rate limiting and throttling
- Input validation and sanitization

Data Protection

- Encryption at rest and in transit
- Key management systems
- HTTPS/TLS best practices
- SQL injection and XSS prevention

Monitoring & Observability

Three Pillars of Observability

- Metrics: Business and system metrics (Prometheus, Grafana)
- Logging: Structured logging, log aggregation (ELK stack)
- Tracing: Distributed tracing (Jaeger, Zipkin)

SLAs, SLOs, and Error Budgets

- Service Level Agreements definition
- Error budget calculation and management
- Alert fatigue prevention
- Incident response procedures

2. Advanced Topics for 2026

AI/ML Integration

Real-time ML Inference

- Model serving architectures (TensorFlow Serving, MLflow)
- A/B testing frameworks for ML models
- Feature stores and data pipelines
- Model versioning and rollback strategies

Recommendation Systems

- Collaborative filtering vs content-based filtering
- Real-time vs batch recommendation generation
- Cold start problems and solutions
- Scalability challenges with millions of users

Edge Computing

CDN Optimization

- Dynamic content caching strategies
- Edge computing with Lambda@Edge, Cloudflare Workers
- Global traffic routing and failover
- Cost optimization techniques

Real-time Systems

WebSocket Architecture

- Connection management and scaling
- Message broadcasting patterns
- Fallback mechanisms (long polling, Server-Sent Events)
- Real-time collaboration (Operational Transformation, CRDTs)

Data Engineering

Streaming Data Pipelines

- Lambda vs Kappa architecture
- Stream processing frameworks (Apache Flink, Spark Streaming)
- Data quality and schema evolution
- Real-time analytics and aggregations

3. Common System Design Questions by Company

Google (L4-L6)

Focus Areas: Distributed systems, large-scale data processing

- Design Google Search (L5-L6)
- Design YouTube video streaming (L4-L5)
- Design Google Maps/Navigation (L5-L6)
- Design Google Drive file storage (L4-L5)
- Design a distributed cache system (L4)

Meta (E4-E6)

Focus Areas: Social media scale, real-time features

- Design Facebook News Feed (E5-E6)
- Design Instagram Stories (E4-E5)
- Design WhatsApp messaging (E5-E6)
- Design Facebook Live streaming (E5-E6)
- Design a social media analytics system (E4)

Amazon (SDE II-Principal)

Focus Areas: E-commerce, high availability, cost optimization

- Design Amazon product catalog (SDE II)
- Design Amazon Prime Video (Senior SDE)
- Design Amazon payment system (Senior SDE)
- Design AWS S3-like storage (Principal)
- Design recommendation engine (SDE II-Senior)

Microsoft (L60-L65)

Focus Areas: Enterprise solutions, hybrid cloud

- Design Microsoft Teams (L62-L63)
- Design Office 365 collaboration (L63-L64)
- Design Azure Blob Storage (L62-L63)
- Design Outlook email system (L61-L62)
- Design Xbox Live gaming platform (L63-L64)

Netflix (Senior-Staff)

Focus Areas: Content delivery, recommendation systems

- Design Netflix video streaming (Senior-Staff)
- Design Netflix recommendation system (Senior-Staff)
- Design content encoding pipeline (Senior)
- Design A/B testing platform (Senior-Staff)

Uber (L4-L6)

Focus Areas: Real-time systems, geolocation services

- Design Uber ride matching (L5-L6)
- Design Uber Eats delivery (L4-L5)
- Design real-time location tracking (L5)
- Design surge pricing system (L5-L6)

4. Step-by-Step Interview Approach

Phase 1: Requirements Gathering (5-7 minutes)

Functional Requirements

- Core features and user workflows
- User types and their interactions
- Success metrics and business goals

Non-Functional Requirements

- Scale: Daily/Monthly active users
- Performance: Latency and throughput requirements
- Availability: Uptime expectations
- Consistency: Data consistency requirements

Example Questions to Ask:

- "How many users are we expecting?"
- "What's the read-to-write ratio?"
- "Do we need real-time updates?"
- "What are the geographical distribution requirements?"

Phase 2: Capacity Estimation (5-8 minutes)

User Calculations

Daily Active Users (DAU): 100M

Average requests per user: 50/day

Peak traffic multiplier: 3x

$\text{QPS} = (100\text{M} \times 50) / (24 \times 3600) \times 3 = 173,611 \text{ QPS}$

Storage Calculations

Photos per user per day: 2

Average photo size: 2MB

$\text{Storage per day} = 100\text{M} \times 2 \times 2\text{MB} = 400\text{TB/day}$

$\text{Storage per year} = 400\text{TB} \times 365 = 146\text{PB/year}$

Bandwidth Calculations

$\text{Upload bandwidth} = 173,611 \text{ QPS} \times 2\text{MB} = 347\text{GB/s}$

$\text{Download bandwidth (assuming 10:1 read ratio)} = 3.47\text{TB/s}$

Phase 3: High-Level Design (10-12 minutes)

Core Components

- Load balancers
- Web servers
- Application servers
- Databases
- Caches
- CDNs
- Message queues

Data Flow Diagram

User → Load Balancer → Web Server → App Server → Database



Cache/CDN

Phase 4: Detailed Design (15-20 minutes)

Database Schema Design

sql

Users Table:

- user_id (Primary Key)
- username, email, created_at
- profile_data (JSON)

Posts Table:

- post_id (Primary Key)
- user_id (Foreign Key)
- content, media_urls
- created_at, updated_at

API Design

POST /api/v1/posts

GET /api/v1/posts/{post_id}

GET /api/v1/users/{user_id}/feed?limit=20&cursor=xyz

PUT /api/v1/posts/{post_id}

DELETE /api/v1/posts/{post_id}

Phase 5: Scale & Optimize (8-10 minutes)

Bottleneck Identification

- Database becomes read/write bottleneck
- Single points of failure
- Cache invalidation complexity
- Network bandwidth limitations

Scaling Solutions

- Database sharding and replication
- Microservices decomposition
- Asynchronous processing with queues
- CDN for static content delivery

Phase 6: Trade-offs Discussion (5-8 minutes)

Common Trade-offs

- Consistency vs Availability (CAP theorem)
- Cost vs Performance
- Complexity vs Maintainability
- Security vs Usability

5. Hands-On Practice Plan (16-Week Schedule)

Weeks 1-2: Fundamentals

- **Study:** Scalability principles, CAP theorem
- **Practice:** Design a simple URL shortener
- **Exercise:** Calculate capacity for 1M users system

Weeks 3-4: Database Systems

- **Study:** SQL vs NoSQL, sharding strategies
- **Practice:** Design database schema for social media
- **Exercise:** Implement consistent hashing algorithm

Weeks 5-6: Caching & Load Balancing

- **Study:** Multi-level caching, LB algorithms
- **Practice:** Design a distributed cache system
- **Exercise:** Implement LRU cache with TTL

Weeks 7-8: Message Queues & APIs

- **Study:** Kafka, event-driven architecture
- **Practice:** Design a chat application
- **Exercise:** Build REST API with rate limiting

Weeks 9-10: Advanced Topics

- **Study:** Microservices, event sourcing
- **Practice:** Design Netflix-like streaming service
- **Exercise:** Implement circuit breaker pattern

Weeks 11-12: Real-time Systems

- **Study:** WebSockets, real-time data processing

- **Practice:** Design Uber ride-matching system
- **Exercise:** Build real-time notification system

Weeks 13-14: AI/ML Integration

- **Study:** Recommendation systems, ML serving
- **Practice:** Design recommendation engine
- **Exercise:** Implement collaborative filtering

Weeks 15-16: Mock Interviews

- **Practice:** 8-10 complete system design interviews
- **Focus:** Time management and communication
- **Review:** Common mistakes and improvements

Daily Practice Structure (2-3 hours)

- **30 minutes:** Concept review and reading
- **60 minutes:** System design practice
- **30 minutes:** Coding system components
- **30 minutes:** Diagram drawing and presentation

6. Essential Resources & Tools

Books

1. **"Designing Data-Intensive Applications"** by Martin Kleppmann
 - Deep dive into distributed systems concepts
 - Real-world examples and trade-offs
2. **"System Design Interview - An Insider's Guide"** by Alex Xu
 - Interview-focused approach
 - Step-by-step problem solving
3. **"Building Microservices"** by Sam Newman
 - Microservices architecture patterns
 - Practical implementation guidance

Online Platforms

1. **Educative.io**
 - "Grokking the System Design Interview"
 - Interactive learning with diagrams

2. LeetCode Premium

- System design questions with discussions
- Company-specific problem sets

3. High Scalability Blog

- Real-world architecture case studies
- Industry best practices

Practice Tools

1. Draw.io / Lucidchart

- Architecture diagram creation
- Collaborative editing features

2. Excalidraw

- Quick sketching and wireframing
- Easy sharing and collaboration

3. Figma

- UI/UX design for system interfaces
- Component libraries

Mock Interview Platforms

1. **Pramp** - Free peer-to-peer interviews
2. **Exponent** - Tech interview preparation
3. **InterviewBit** - Structured practice sessions
4. **Interviewing.io** - Anonymous practice with engineers

7. Real-World Case Studies

Netflix Streaming Architecture

Key Components:

- **Content Delivery Network (CDN):** 95% of traffic served from edge locations
- **Microservices:** 700+ microservices for different functionalities
- **Chaos Engineering:** Proactive failure testing with Chaos Monkey
- **A/B Testing:** Continuous experimentation platform

Scale Numbers:

- 230M+ global subscribers
- 15,000+ titles in catalog

- 1 billion hours watched per week
- 200+ countries served

Architecture Decisions:

- AWS for cloud infrastructure (no data centers)
- Cassandra for user viewing history (write-heavy workload)
- Elasticsearch for search and recommendations
- Kafka for real-time event streaming

Uber's Real-time Location Tracking

Key Components:

- **DISCO (Distributed Compute):** Real-time dispatch system
- **Ringpop:** Consistent hashing for load distribution
- **Geospatial Indexing:** QuadTree for location queries
- **Supply-Demand Forecasting:** ML models for driver positioning

Scale Numbers:

- 93M monthly active users
- 18M trips per day
- 5M drivers globally
- Sub-second matching requirements

Technical Challenges:

- GPS accuracy and map-matching algorithms
- Real-time ETAs with traffic data
- Dynamic pricing algorithms
- Cross-border trip handling

WhatsApp Messaging System

Key Components:

- **XMPP Protocol:** Modified for mobile optimization
- **Erlang/OTP:** Actor model for concurrent connections
- **FreeBSD Servers:** Custom kernel optimizations
- **End-to-End Encryption:** Signal Protocol implementation

Scale Numbers:

- 2B+ users worldwide
- 100B messages per day
- 50 engineers maintaining the system
- 99.9%+ uptime requirement

Architecture Decisions:

- Single-threaded server design for simplicity
- Message queuing for offline users
- Minimal metadata storage for privacy
- Aggressive caching for contact lists

Instagram Photo Sharing

Key Components:

- **Django Framework:** Python web application
- **PostgreSQL:** User data and relationships
- **Cassandra:** Photo metadata and activity feeds
- **Amazon S3:** Photo storage with CDN

Scale Numbers:

- 1B+ monthly active users
- 95M photos uploaded daily
- 4.2B likes per day
- 500M+ Instagram Stories daily

Technical Decisions:

- Asynchronous task processing with Celery
- Push vs pull feed generation strategies
- Image resizing and optimization pipeline
- Redis for real-time features (Stories, Live)

Twitter/X Timeline Generation

Key Components:

- **Fanout Service:** Timeline computation and caching
- **Tweet Service:** 140/280 character content management

- **Social Graph:** Following/follower relationships
- **Trend Detection:** Real-time hashtag analytics

Scale Numbers:

- 450M+ monthly active users
- 500M tweets per day
- 350,000 tweets per minute during peak
- 40+ languages supported

Technical Challenges:

- Celebrity user fanout optimization
- Real-time trend detection algorithms
- Spam and abuse detection systems
- Global content replication and caching

8. Common Pitfalls & How to Avoid Them

Over-engineering Solutions

Pitfall: Immediately jumping to complex architectures **Solution:** Start simple, then scale based on requirements **Example:** Don't use microservices for a system with 10K users

Ignoring Trade-offs

Pitfall: Not discussing pros/cons of design decisions **Solution:** Always mention alternatives and explain choices **Example:** "We chose NoSQL for scalability, but we lose ACID guarantees"

Poor Time Management

Pitfall: Spending too much time on one section **Solution:** Follow the structured timeline religiously

Recommended Timing:

- Requirements: 5-7 minutes
- Estimation: 5-8 minutes
- High-level: 10-12 minutes
- Detailed: 15-20 minutes
- Scaling: 8-10 minutes
- Trade-offs: 5-8 minutes

Inadequate Requirements Clarification

Pitfall: Making assumptions without asking questions **Solution:** Always clarify ambiguous requirements

Key Questions:

- "Are we designing for mobile, web, or both?"
- "What's more important: consistency or availability?"
- "Do we need to support offline functionality?"

Missing Monitoring and Reliability

Pitfall: Focusing only on happy path scenarios **Solution:** Always discuss observability and failure scenarios **Include:** Logging, monitoring, alerting, disaster recovery

9. Company-Specific Focus Areas

Google: Distributed Systems Excellence

Key Focus Areas:

- Large-scale data processing (MapReduce, BigTable concepts)
- Global infrastructure and edge computing
- Search and information retrieval algorithms
- Machine learning integration at scale

Preparation Tips:

- Study Google's published papers (GFS, MapReduce, Spanner)
- Focus on consistency models and distributed consensus
- Understand PageRank and search ranking algorithms
- Practice with planetary-scale scenarios (billions of users)

Meta: Social Media Scale

Key Focus Areas:

- Real-time social features (News Feed, messaging)
- Graph-based data structures and algorithms
- Content recommendation and ranking
- Mobile-first architecture considerations

Preparation Tips:

- Study Facebook's architecture evolution papers
- Focus on social graph traversal algorithms
- Understand feed ranking and content delivery

- Practice with real-time features and WebSocket scaling

Amazon: High Availability & Cost Optimization

Key Focus Areas:

- E-commerce transaction processing
- Inventory management systems
- Cost-effective cloud architecture
- High availability and disaster recovery

Preparation Tips:

- Study Amazon's architecture principles (Two-Pizza teams)
- Focus on eventual consistency and CAP theorem
- Understand microservices decomposition strategies
- Practice with cost optimization scenarios

Microsoft: Enterprise Solutions

Key Focus Areas:

- Enterprise software architecture
- Hybrid cloud solutions
- Identity and access management
- Office productivity suite scaling

Preparation Tips:

- Focus on enterprise security and compliance
- Study hybrid cloud architecture patterns
- Understand Active Directory and identity systems
- Practice with B2B software scaling challenges

Netflix: Content Delivery & Personalization

Key Focus Areas:

- Video streaming infrastructure
- Content recommendation algorithms
- A/B testing and experimentation platforms
- Chaos engineering and reliability

Preparation Tips:

- Study Netflix's tech blog extensively
- Focus on CDN and video encoding systems
- Understand recommendation system architectures
- Practice with global content delivery scenarios

Uber: Real-time Systems & Geolocation

Key Focus Areas:

- Real-time matching algorithms
- Geospatial data processing
- Dynamic pricing systems
- Supply and demand forecasting

Preparation Tips:

- Study geospatial indexing algorithms (QuadTree, R-tree)
- Focus on real-time processing and stream computing
- Understand marketplace dynamics and pricing
- Practice with location-based service design

10. 2026-Specific Trends to Know

Serverless Architecture Adoption

Key Concepts:

- Function-as-a-Service (FaaS) patterns
- Event-driven serverless architectures
- Cold start optimization techniques
- Serverless database solutions (Aurora Serverless, DynamoDB On-Demand)

Design Implications:

- Auto-scaling based on event triggers
- Stateless function design principles
- Cost optimization through pay-per-execution
- Integration with managed cloud services

AI-Driven System Optimization

Key Concepts:

- Predictive auto-scaling using ML models
- Intelligent load balancing and traffic routing
- Automated performance tuning and optimization
- AI-powered incident detection and resolution

Implementation Examples:

- ML models predicting traffic spikes for proactive scaling
- AI-driven database query optimization
- Intelligent caching decisions based on access patterns
- Automated A/B testing and feature flag management

Multi-Cloud Strategies

Key Concepts:

- Cloud-agnostic architecture design
- Cross-cloud data replication and sync
- Disaster recovery across cloud providers
- Cost optimization through cloud arbitrage

Technical Challenges:

- Data consistency across different cloud providers
- Network latency and bandwidth costs
- Vendor-specific service dependencies
- Unified monitoring and observability

Sustainability in System Design

Key Concepts:

- Green computing and energy-efficient architectures
- Carbon footprint measurement and optimization
- Sustainable data center design principles
- Energy-aware workload scheduling

Design Considerations:

- Choosing regions with renewable energy sources
- Optimizing resource utilization and reducing waste

- Implementing efficient data compression and storage
- Considering lifecycle costs of hardware and infrastructure

Privacy-First Architecture

Key Concepts:

- Privacy by design principles
- Zero-knowledge architecture patterns
- Differential privacy implementation
- GDPR and CCPA compliance by design

Technical Implementation:

- Data minimization and purpose limitation
- Encrypted data processing and homomorphic encryption
- Decentralized identity and authentication systems
- Privacy-preserving analytics and ML

Quantum-Resistant Security

Key Concepts:

- Post-quantum cryptography algorithms
- Quantum-safe key exchange protocols
- Migration strategies for existing systems
- Hybrid classical-quantum security approaches

Preparation Areas:

- Understanding quantum computing threats to current encryption
- Learning about lattice-based and hash-based cryptography
- Planning for cryptographic agility in system design
- Considering quantum key distribution for high-security applications

Final Preparation Checklist

Technical Knowledge Verification

- ☐ Can explain CAP theorem with real examples
- ☐ Understands database sharding and consistent hashing
- ☐ Can design REST APIs with proper HTTP semantics
- ☐ Knows caching strategies and invalidation patterns

- ☐ Understands load balancing algorithms and trade-offs
- ☐ Can explain microservices vs monolith trade-offs
- ☐ Knows message queue patterns and event-driven architecture
- ☐ Understands monitoring, logging, and observability
- ☐ Can calculate capacity and estimate system resources
- ☐ Knows security best practices and common vulnerabilities

Communication Skills

- ☐ Can ask clarifying questions effectively
- ☐ Explains technical concepts clearly to non-technical audience
- ☐ Discusses trade-offs and alternatives
- ☐ Manages time effectively during mock interviews
- ☐ Draws clear and understandable system diagrams
- ☐ Handles follow-up questions confidently
- ☐ Admits knowledge gaps honestly and suggests alternatives

Practice Milestones

- ☐ Completed 20+ system design problems
- ☐ Conducted 10+ mock interviews with peers
- ☐ Drew 50+ system architecture diagrams
- ☐ Practiced capacity estimation for various scales
- ☐ Reviewed 10+ real-world architecture case studies
- ☐ Implemented 5+ system design components in code
- ☐ Studied company-specific architecture blogs and papers

Company Research

- ☐ Read engineering blogs of target companies
- ☐ Studied recent architecture decisions and migrations
- ☐ Understood company-specific scale and challenges
- ☐ Researched team structures and engineering culture
- ☐ Reviewed recent job postings for technical requirements
- ☐ Connected with current employees on LinkedIn
- ☐ Prepared company-specific questions to ask interviewers

Remember: System design interviews are as much about communication and thought process as they are about technical knowledge. Practice explaining your reasoning clearly and always consider the business context and trade-offs in your solutions.

Good luck with your interviews!